Reinforcement Learning

Other topics

Mario Martin

CS-UPC

May 24, 2018

Topics not discussed

- There are still some hot topics we haven't covered
  - Model based Reinforcement Learning
  - Inverse Reinforcement Learning (IRL)
  - Partial Observability: Memory approaches
  - Multi Agent Reinforcement Learning (MARL)
  - Hierarchical reinforcement learning
  - Exploration vs. Explotation
  - Transfer Learning
  - Robotics

- Today, we'll mention some of them

Model based Reinforcement Learning

- Model-free methods do not need transition probabilities ($P^{a'}_{s,a}$)
- In case where we have these transitions we can do more things:
  - Planning: You can learn with complete backups using probabilities instead of samples. Of with different degrees of depth (n-steps)
  - Rollouts: You can learn generating predicted trials
  - Monte-Carlo Tree search: Combination of two previous cases.
  - Smarter exploration
- In case you don't have the model, learn it explicitly from samples and use it as it were given.

Inverse Reinforcement Learning

- Consists in, given an optimal policy (or examples of the policy), obtain the reward function.
- In some cases we cannot apply RL because we reinforcement function is very complex
- But we have examples of the policy we want to learn
- In these cases, IRL allows to get the reward function and from that learn the policy
- More robust than learning from examples

Partial Observability

- In a lot of cases the agent has not complete information of the state.
- The problem is not anymore an MDP.
- How to solve these case?
  1. Formalize as a POMDP: MDP extended with set of observations O and probability of each observation given the true state. Agent work with a *belief vector* of probabilities of being in each state. Solve with dedicated algorithms
  2. Works with memory as a way tot disambiguate the state. Simple approaches like window of last n perceptions, or more interesting ones using LSTM
  3. Find an stochastic policy using policy search methods or evolutionary methods

Hierarchic RL

- In some cases a complex task can be decomposed in simpler tasks.
- Learning is simplified when first these tasks are learnt.
- Several ways to find that:
  1. Using subrewards for subactions (reward shaping)
  2. Discover them authomatically
- Useful for transfer learning

Transfer RL

- Can we extend knowledge generated in one task to a different task?

Multi-agent RL

- All cases we have seen assume the agent is the only one that executes actions in the environment
- In cases where there are also other agents, can we learn?
- Use of game theory and assumptions about the other agents
- Depending on the goals of the agent, we have cooperative or competitive learning
- Two-players games are an special case.
  - Backgammon: Neurogammon (Tesauro 1994)
  - Go: Alpha-go (Silver et al. 2016) and Alpha-go Zero (Silver et al. 2017)
  - Chess: AlphaZero (Silver et al. 2017)